

# ALL\_Models

July 1, 2024

## 1 Model Implementation

This code covers of implementing Polynomial Regression, Decision Tree and LSTM

```
[1]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import PolynomialFeatures
      from sklearn.linear_model import LinearRegression
      from sklearn.tree import DecisionTreeRegressor
      from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
      import numpy as np
      import matplotlib.pyplot as plt
```

### 1.0.1 Data Loading and Initial Exploration

In this section, we load the S&P 500 index data from a CSV file and display the first and last few rows to understand the dataset's structure.

```
[2]: # import yfinance as yf
      # df = yf.download('^SPX', start = '1995-12-27')

      df = pd.read_csv('SPX1995.csv')
```

```
[3]: df.head()
```

```
[3]:      Date      Open      High      Low      Close  Adj Close  \
0  1995-01-03  459.209991  459.269989  457.200012  459.109985  459.109985
1  1995-01-04  459.130005  460.720001  457.559998  460.709991  460.709991
2  1995-01-05  460.730011  461.299988  459.750000  460.339996  460.339996
3  1995-01-06  460.380005  462.489990  459.470001  460.679993  460.679993
4  1995-01-09  460.670013  461.769989  459.739990  460.829987  460.829987

      Volume
0  262450000
1  319510000
2  309050000
3  308070000
4  278790000
```

```
[4]: df.tail()
```

```
[4]:
```

	Date	Open	High	Low	Close \
7336	2024-02-23	5100.919922	5111.060059	5081.459961	5088.799805
7337	2024-02-26	5093.000000	5097.660156	5068.910156	5069.529785
7338	2024-02-27	5074.600098	5080.689941	5057.290039	5078.180176
7339	2024-02-28	5067.200195	5077.370117	5058.350098	5069.759766
7340	2024-02-29	5085.359863	5104.990234	5061.890137	5096.270020

	Adj Close	Volume
7336	5088.799805	3672790000
7337	5069.529785	3683930000
7338	5078.180176	3925950000
7339	5069.759766	3789370000
7340	5096.270020	5219740000

```
[5]: df.shape
```

```
[5]: (7341, 7)
```

```
[ ]:
```

### 1.0.2 Data Cleaning: Handling Missing Values and Duplicates

This section focuses on identifying and addressing any missing or duplicated data entries to ensure the quality and reliability of the dataset for further analysis.

```
[6]: missing_values = df.isnull().sum()
df_duplicated= df.duplicated().sum().any()

# here we drop rows if there is missing values
df_cleaned = df.dropna()

print("Missing values in each column:\n", missing_values)
print("\n \n duplicated values : ", df_duplicated)
```

Missing values in each column:

```
Date          0
Open          0
High          0
Low           0
Close         0
Adj Close     0
Volume        0
dtype: int64
```

```
duplicated values : False
```

### 1.0.3 Column Removal

In this section, we remove columns from the dataset that are not needed for our analysis.

```
[7]: columns_to_drop = ['Adj Close']
df = df.drop(columns_to_drop, axis=1)
df
```

```
[7]:
```

	Date	Open	High	Low	Close \
0	1995-01-03	459.209991	459.269989	457.200012	459.109985
1	1995-01-04	459.130005	460.720001	457.559998	460.709991
2	1995-01-05	460.730011	461.299988	459.750000	460.339996
3	1995-01-06	460.380005	462.489990	459.470001	460.679993
4	1995-01-09	460.670013	461.769989	459.739990	460.829987
...	...	...	...	...	...
7336	2024-02-23	5100.919922	5111.060059	5081.459961	5088.799805
7337	2024-02-26	5093.000000	5097.660156	5068.910156	5069.529785
7338	2024-02-27	5074.600098	5080.689941	5057.290039	5078.180176
7339	2024-02-28	5067.200195	5077.370117	5058.350098	5069.759766
7340	2024-02-29	5085.359863	5104.990234	5061.890137	5096.270020
...	...	...	...	...	...
7336	2024-02-23	3672790000	3683930000	3925950000	3789370000
7337	2024-02-26	3683930000	3925950000	3789370000	5219740000
7338	2024-02-27	3925950000	3789370000	5219740000	
7339	2024-02-28	3789370000	5219740000		
7340	2024-02-29	5219740000			

[7341 rows x 6 columns]

### 1.0.4 Visualization of S&P 500 Stock Prices

In this section, we convert the 'Date' column to datetime format for proper indexing and plot the S&P 500 closing and opening prices over time to visualize trends and patterns in the data.

```
[8]: import matplotlib.pyplot as plt
import seaborn as sns

# Set the style of seaborn
sns.set(style='darkgrid')

# Convert 'Date' to datetime
```

```

df['Date'] = pd.to_datetime(df['Date'])


# Plotting the closing prices against the date
plt.figure(figsize=(14, 7))
plt.plot(df['Date'], df['Close'], label='Close Price')

plt.plot(df['Date'], df['Open'], label='Open Price', alpha=0.5)

# Labels and Title
plt.xlabel('Date')
plt.ylabel('Price')
plt.title('S&P 500 Stock Prices')
plt.legend()

# Show plot
plt.show()

```



### 1.0.5 Feature Preparation

Preparing the dataset for modeling by selecting the 'Open' and 'Volume' as features and 'Close' as the target variable. Converting 'Date' to a numerical format for use in polynomial features.

```

[9]: features = [ 'Open', 'High', 'Low', 'Volume']
    target = 'Close'

# We split the data into features and target

```

```
X = df[features]
y = df[target]
```

## 1.1 Polynomial Regression Implementation

### 1.1.1 Data Splitting For Polynomial Regression

Dividing the data into training and test sets to validate the performance of our model. Ensuring a fair distribution without shuffling due to the time-series nature of the data.

```
[10]: degree = 2
poly_features = PolynomialFeatures(degree=degree)

X_poly = poly_features.fit_transform(X)

# we save the indexes before the split
original_indexes = df.index

# here we split the dataset into training and testing sets by making 80% for
↳ training and 20% for test
X_train, X_test, y_train, y_test = train_test_split(X_poly, df['Close'],
↳ test_size=0.2, shuffle=False)

print("the train data: --> ", X_train.shape)
print("the test data: --> ", X_test.shape)
```

```
the train data: --> (5872, 15)
the test data: --> (1469, 15)
```

### 1.1.2 Polynomial Model Initialization and Fitting

Initializing the Linear Regression model and fitting it to the polynomial-transformed training data to capture non-linear patterns in stock prices.

```
[11]: polynomial_model = LinearRegression()
polynomial_model.fit(X_train, y_train)
```

```
[11]: LinearRegression()
```

### 1.1.3 Preparing Index Alignment Before Making Predictions

Before making predictions, it's important to retrieve and store the original indexes of our training and testing sets. This allows us to maintain a reference to the original dates of our data points, ensuring that after we make our predictions, we can accurately analyze and visualize the results in the context of their specific times in the dataset.

```
[12]: # Retrieve the original indexes for train and test sets
train_indexes = original_indexes[:len(y_train)]
test_indexes = original_indexes[len(y_train):]
```

```
print(f'Training data index range: {train_indexes.min()} to {train_indexes.
↪max()}')
print(f'Testing data index range: {test_indexes.min()} to {test_indexes.max()}')
```

Training data index range: 0 to 5871

Testing data index range: 5872 to 7340

### 1.1.4 Model Predictions

After training our regression model, we proceed to make predictions on both the training and testing datasets. These predictions will allow us to evaluate the model's performance by comparing the predicted stock prices against the actual closing prices. It's crucial to ensure that the predictions align with the original data's timeline, hence the index retrieval before this step.

```
[13]: # Predict on the training set for visualization purposes
y_train_poly_pred = polynomial_model.predict(X_train)

# Make predictions on the testing data
y_test_poly_pred = polynomial_model.predict(X_test)
```

### 1.1.5 Organizing and Inspecting Polynomial Regression Prediction Results in Table Form

In this section, we consolidate the predictions with the actual values into structured DataFrames, aligning them with their corresponding dates. This organization is essential for an intuitive inspection of the model's predictive accuracy. It also lays the groundwork for subsequent analysis, such as calculating error metrics and visualizing the results.

```
[14]: if 'Date' in df.columns:
        df['Date'] = pd.to_datetime(df['Date'])

train_dates = df['Date'].iloc[:len(y_train)].reset_index(drop=True)
test_dates = df['Date'].iloc[-len(y_test):].reset_index(drop=True)

poly_train_results = pd.DataFrame({
    'Date': train_dates,
    'Actual_Close': y_train.reset_index(drop=True),
    'Poly_Predicted_Close': y_train_poly_pred
})

poly_test_results = pd.DataFrame({
    'Date': test_dates,
    'Actual_Close': y_test.reset_index(drop=True),
    'Poly_Predicted_Close': y_test_poly_pred
})
```

```
print("poly_train_results \n")
print( poly_train_results , "\n \n")

print("poly_test_results \n")
print( poly_test_results)
```

poly\_train\_results

	Date	Actual_Close	Poly_Predicted_Close
0	1995-01-03	459.109985	667.145371
1	1995-01-04	460.709991	675.250452
2	1995-01-05	460.339996	673.398711
3	1995-01-06	460.679993	674.563202
4	1995-01-09	460.829987	670.228492
...	...	...	...
5867	2018-04-23	2670.290039	2913.400999
5868	2018-04-24	2634.560059	2925.766888
5869	2018-04-25	2639.399902	2846.865873
5870	2018-04-26	2666.939941	2912.266104
5871	2018-04-27	2669.909912	2878.095985

[5872 rows x 3 columns]

poly\_test\_results

	Date	Actual_Close	Poly_Predicted_Close
0	2018-04-30	2648.050049	2880.951228
1	2018-05-01	2654.800049	2862.861414
2	2018-05-02	2635.669922	2834.173511
3	2018-05-03	2629.729980	2827.491682
4	2018-05-04	2663.419922	2999.733193
...	...	...	...
1464	2024-02-23	5088.799805	8617.675638
1465	2024-02-26	5069.529785	8548.768183
1466	2024-02-27	5078.180176	8455.260613
1467	2024-02-28	5069.759766	8470.372507
1468	2024-02-29	5096.270020	8424.977497

[1469 rows x 3 columns]

### 1.1.6 Performance Metrics Evaluation for Polynomial Regression

In this segment, we compute and display the performance metrics for both the training and testing datasets. This evaluation involves Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE) and Root Mean Squared Error (RMSE). These metrics help to quantify the accuracy of our model and reveal how well the predictions match up with the actual stock prices.

```
[15]: from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

# Function to calculate MAPE
def mean_absolute_percentage_error(y_true, y_pred):
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

# Calculate metrics for the training set
poly_train_mae = mean_absolute_error(poly_train_results['Actual_Close'],
    ↪poly_train_results['Poly_Predicted_Close'])
poly_train_rmse = np.
    ↪sqrt(mean_squared_error(poly_train_results['Actual_Close'],
    ↪poly_train_results['Poly_Predicted_Close']))
poly_train_mape =
    ↪mean_absolute_percentage_error(poly_train_results['Actual_Close'],
    ↪poly_train_results['Poly_Predicted_Close'])

# Calculate metrics for the testing set
poly_test_mae = mean_absolute_error(poly_test_results['Actual_Close'],
    ↪poly_test_results['Poly_Predicted_Close'])
poly_test_rmse = np.sqrt(mean_squared_error(poly_test_results['Actual_Close'],
    ↪poly_test_results['Poly_Predicted_Close']))
poly_test_mape =
    ↪mean_absolute_percentage_error(poly_test_results['Actual_Close'],
    ↪poly_test_results['Poly_Predicted_Close'])

# Print out the metrics for the training set
print("Training set metrics:")
print(f"Mean Absolute Error (MAE): {poly_train_mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {poly_train_rmse:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {poly_train_mape:.2f}%")

# Print out the metrics for the testing set
print("\nTesting set metrics:")
print(f"Mean Absolute Error (MAE): {poly_test_mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {poly_test_rmse:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {poly_test_mape:.2f}%")
```

Training set metrics:  
Mean Absolute Error (MAE): 61.77  
Root Mean Squared Error (RMSE): 81.99  
Mean Absolute Percentage Error (MAPE): 5.77%

Testing set metrics:  
Mean Absolute Error (MAE): 1301.87  
Root Mean Squared Error (RMSE): 1576.16  
Mean Absolute Percentage Error (MAPE): 32.09%



[ ]:

### 1.1.7 Visualization of Model Predictions Against Actual Data

In accordance with our project's aim to assess machine learning model efficacy, this visualization plots predicted stock prices from our model against the actual S&P 500 closing prices. The graph provides a visual representation of the model's performance over time, showcasing the alignment of predictions with real-world data. This step is crucial for a comprehensive evaluation, allowing for a clear, intuitive understanding of the model's predictive capabilities in both training and testing phases.

```
[16]: import matplotlib.pyplot as plt
import pandas as pd

# here we Combine the train and the test results into a single DataFrame
poly_combined_results = pd.concat([poly_train_results, poly_test_results])

# Here we Convert 'Date' to datetime and then we sort by date to ensure correct
    ↪plotting order
poly_combined_results['Date'] = pd.to_datetime(poly_combined_results['Date'])
poly_combined_results.sort_values('Date', inplace=True)

# Set 'Date' as the index for plotting
poly_combined_results.set_index('Date', inplace=True)

# Plot the actual close prices
plt.figure(figsize=(14,8))
plt.plot(poly_combined_results['Actual_Close'], label='Actual Close',
    ↪color='blue')


# Plot the training predictions
plt.plot(poly_train_results['Date'],
    ↪poly_train_results['Poly_Predicted_Close'], label='Training Predictions',
    ↪color='orange', linestyle='--')

# Plot the testing predictions
plt.plot(poly_test_results['Date'], poly_test_results['Poly_Predicted_Close'],
    ↪label='Testing Predictions', color='green', linestyle='--')

plt.xticks(fontsize=14) # Larger font size for the x-axis ticks
plt.yticks(fontsize=14) # Larger font size for the y-axis ticks

# Here we just added the labels and title for the graph
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price', fontsize=18)
plt.title('Actual Close vs Predictions', fontsize=18)
plt.legend(prop={'size': 18, 'weight': 'bold'})
```

```
plt.savefig('poly_result_graph.pdf', format='pdf')
plt.show()
```



## 1.2 Decision Tree Implementation

### 1.2.1 Data Splitting For Decision Tree

Dividing the data into training and test sets to validate the performance of our model. Ensuring a fair distribution without shuffling due to the time-series nature of the data.

```
[17]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42, shuffle=False)
```

### 1.2.2 Initializing and Fitting the Decision Tree Regressor

Initializing the Decision Tree Regressor with specified parameters and fitting it to the training data. This step involves configuring the model with constraints to prevent overfitting and training it to learn patterns from the data

```
[18]: # Initializing the Decision Tree Regressor with specified parameters
decision_model = DecisionTreeRegressor(max_depth=10, min_samples_split=50,
↳ min_samples_leaf=20, random_state=42)

# Fitting the Decision Tree model to the training data
decision_model.fit(X_train, y_train)
```

```
[18]: DecisionTreeRegressor(max_depth=10, min_samples_leaf=20, min_samples_split=50,
                             random_state=42)
```

### 1.2.3 Preparing Index Alignment Before Making Predictions

Before making predictions, it's important to retrieve and store the original indexes of our training and testing sets. This allows us to maintain a reference to the original dates of our data points, ensuring that after we make our predictions, we can accurately analyze and visualize the results in the context of their specific times in the dataset.

```
[19]: # Retrieve the original indexes for train and test sets
train_indexes = original_indexes[:len(y_train)]
test_indexes = original_indexes[len(y_train):]

# You can now use train_indexes and test_indexes as they contain the original_
↳ DataFrame indexes
print(f'Training data index range: {train_indexes.min()} to {train_indexes.
↳ max()}')
print(f'Testing data index range: {test_indexes.min()} to {test_indexes.max()}')
```

Training data index range: 0 to 5871

Testing data index range: 5872 to 7340

### 1.2.4 Making Predictions with the Decision Tree Model

After training the Decision Tree model, the next step is to make predictions on both the training and testing datasets. This helps in evaluating the model's performance and understanding how well it generalizes to new, unseen data.

```
[20]: # Predict on the training set for visualization purposes
y_train_decision_pred = decision_model.predict(X_train)

# Make predictions on the testing data
y_test_decision_pred = decision_model.predict(X_test)
```

### 1.2.5 Organizing and Inspecting Decision Tree Prediction Results in Table Form

In this section, we consolidate the predictions with the actual values into structured DataFrames, aligning them with their corresponding dates. This organization is essential for an intuitive inspection of the model's predictive accuracy. It also lays the groundwork for subsequent analysis, such as calculating error metrics and visualizing the results.

```
[21]: if 'Date' in df.columns:
        df['Date'] = pd.to_datetime(df['Date'])

train_dates = df['Date'].iloc[:len(y_train)].reset_index(drop=True)
test_dates = df['Date'].iloc[-len(y_test):].reset_index(drop=True)
```

```

# here we create a DataFrame to store the actual and predicted closing prices,
↳ for the training set along with their dates
decision_train_results = pd.DataFrame({
    'Date': train_dates,
    'Actual_Close': y_train.reset_index(drop=True),
    'Predicted_Close': y_train_decision_pred
})

# Create a DataFrame to store the actual and predicted closing prices for the,
↳ testing set along with their dates
decision_test_results = pd.DataFrame({
    'Date': test_dates,
    'Actual_Close': y_test.reset_index(drop=True),
    'Predicted_Close': y_test_decision_pred
})

print("decision_train_results \n")
print( decision_train_results , "\n \n")

print("decision_test_results \n")
print( decision_test_results)

```

decision\_train\_results

	Date	Actual_Close	Predicted_Close
0	1995-01-03	459.109985	465.801739
1	1995-01-04	460.709991	465.801739
2	1995-01-05	460.339996	465.801739
3	1995-01-06	460.679993	465.801739
4	1995-01-09	460.829987	465.801739
...	...	...	...
5867	2018-04-23	2670.290039	2650.707147
5868	2018-04-24	2634.560059	2650.707147
5869	2018-04-25	2639.399902	2650.707147
5870	2018-04-26	2666.939941	2650.707147
5871	2018-04-27	2669.909912	2650.707147

[5872 rows x 3 columns]

decision\_test\_results

	Date	Actual_Close	Predicted_Close
0	2018-04-30	2648.050049	2650.707147
1	2018-05-01	2654.800049	2650.707147

2	2018-05-02	2635.669922	2650.707147
3	2018-05-03	2629.729980	2650.707147
4	2018-05-04	2663.419922	2650.707147
...	...	...	...
1464	2024-02-23	5088.799805	2808.062512
1465	2024-02-26	5069.529785	2808.062512
1466	2024-02-27	5078.180176	2808.062512
1467	2024-02-28	5069.759766	2808.062512
1468	2024-02-29	5096.270020	2808.062512

[1469 rows x 3 columns]

### 1.2.6 Performance Metrics Evaluation for Decision Tree

In this segment, we compute and display the performance metrics for both the training and testing datasets. This evaluation involves Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE) and Root Mean Squared Error (RMSE). These metrics help to quantify the accuracy of our model and reveal how well the predictions match up with the actual stock prices.

```
[22]: from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

# Function to calculate MAPE
def mean_absolute_percentage_error(y_true, y_pred):
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

# Calculate metrics for the training set
dt_train_mae = mean_absolute_error(decision_train_results['Actual_Close'],
    ↪decision_train_results['Predicted_Close'])
dt_train_rmse = np.
    ↪sqrt(mean_squared_error(decision_train_results['Actual_Close'],
    ↪decision_train_results['Predicted_Close']))
dt_train_mape =
    ↪mean_absolute_percentage_error(decision_train_results['Actual_Close'],
    ↪decision_train_results['Predicted_Close'])

# Calculate metrics for the testing set
dt_test_mae = mean_absolute_error(decision_test_results['Actual_Close'],
    ↪decision_test_results['Predicted_Close'])
dt_test_rmse = np.
    ↪sqrt(mean_squared_error(decision_test_results['Actual_Close'],
    ↪decision_test_results['Predicted_Close']))
dt_test_mape =
    ↪mean_absolute_percentage_error(decision_test_results['Actual_Close'],
    ↪decision_test_results['Predicted_Close'])

# Print out the metrics for the training set
```

```

print("Training set metrics:")
print(f"Mean Absolute Error (MAE): {dt_train_mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {dt_train_rmse:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {dt_train_mape:.2f}%")

# Print out the metrics for the testing set
print("\nTesting set metrics:")
print(f"Mean Absolute Error (MAE): {dt_test_mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {dt_test_rmse:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {dt_test_mape:.2f}%")

```

Training set metrics:  
Mean Absolute Error (MAE): 6.20  
Root Mean Squared Error (RMSE): 8.81  
Mean Absolute Percentage Error (MAPE): 0.49%

Testing set metrics:  
Mean Absolute Error (MAE): 884.28  
Root Mean Squared Error (RMSE): 1111.20  
Mean Absolute Percentage Error (MAPE): 21.32%

### 1.2.7 Visualizing Decision Tree Model Predictions Against Actual Data

In this section, we visualize the predicted stock prices from the Decision Tree model against the actual S&P 500 closing prices. This step provides a visual representation of the model's performance over time, showcasing how well the predictions align with real-world data.

```

[23]: import matplotlib.pyplot as plt
import pandas as pd

# Here we combine train and test results into a single DataFrame.
# This makes it easier to plot and compare the entire dataset's performance in
    ↪ one figure.
decision_combined_results = pd.concat([decision_train_results,
    ↪ decision_test_results])

# Here then we convert 'Date' to datetime and we sort by date to ensure correct
    ↪ plotting order
decision_combined_results['Date'] = pd.
    ↪ to_datetime(decision_combined_results['Date'])
decision_combined_results.sort_values('Date', inplace=True)

# Set 'Date' as the index for plotting
decision_combined_results.set_index('Date', inplace=True)

# Plot the actual close prices
plt.figure(figsize=(14,9))

```

```

plt.plot(decision_combined_results['Actual_Close'], label='Actual Close',
        color='blue')


# Plot the training predictions
plt.plot(decision_train_results['Date'],
        decision_train_results['Predicted_Close'], label='Training Predictions',
        color='orange', linestyle='--')

# Plot the testing predictions
plt.plot(decision_test_results['Date'],
        decision_test_results['Predicted_Close'], label='Testing Predictions',
        color='green', linestyle='--')

# Labels and title are added to provide context to the axes and the plot,
# making it easier to understand.
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

# Add labels and title
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price', fontsize=18)
plt.title('Actual Close vs Predictions', fontsize=18)
plt.legend(prop={'size': 18, 'weight': 'bold'})
# plt.savefig('DT_result_graph.pdf', format='pdf')
plt.savefig('analys_dt_graph.pdf', format='pdf')
plt.show()

```



## 1.3 LSTM Implementation

### 1.3.1 Normalizing Data for LSTM

In this section, we normalize the 'Close' price data using the `MinMaxScaler` from `scikit-learn`. Normalization is a crucial step in preparing data for LSTM (Long Short-Term Memory) networks, as it scales the data to a range suitable for training the neural network, ensuring faster convergence and better performance.

```
[24]: from sklearn.preprocessing import MinMaxScaler

# Here we extract the 'Close' price column and reshape it to a 2D array
data_to_normalize = df['Close'].values.reshape(-1, 1)

# And here we create the MinMaxScaler with the feature range between 0 and 1
# This scaler will transform the data to fit within the specified range
scaler = MinMaxScaler(feature_range=(0, 1))

# We then fits the scaler to the data and transforms it
# This step scales the 'Close' price data to the range [0, 1]
scaled_data = scaler.fit_transform(data_to_normalize)

# Then we display the normalized data
scaled_data
```



```
[24]: array([[0.00000000e+00],
           [3.45040065e-04],
           [2.65250927e-04],
           ...,
           [9.96098939e-01],
           [9.94283084e-01],
           [1.00000000e+00]])
```

### 1.3.2 Creating Sequences for LSTM Model

To prepare the data for the LSTM model, we need to create sequences of a specified length. Each sequence consists of a fixed number of previous time steps, and the output is the next time step. This helps the LSTM model learn the temporal dependencies in the data.

```
[25]: # Function to create sequences and their corresponding labels
# This function generates sequences of 'look_back' length and their
# corresponding labels (next time step).
def create_dataset(dataset, look_back=1):
    dataX, dataY = [], []
    for i in range(len(dataset) - look_back - 1):
        a = dataset[i:(i + look_back), 0]
        dataX.append(a)
        dataY.append(dataset[i + look_back, 0])
    return np.array(dataX), np.array(dataY)

# We used 30 days of previous data points to predict the next value.
look_back = 30

# Create the dataset with sequences, which generates the input-output pairs
# using the create_dataset function.
X, y = create_dataset(scaled_data, look_back)

# LSTM expects input in the form of a 3D array [samples, time steps, features],
# so we reshape accordingly.
X = np.reshape(X, (X.shape[0], look_back, 1))
```

### 1.3.3 Splitting Data into Training and Testing Sets

To train and evaluate the LSTM model effectively, we need to split the data into training and testing sets. The training set is used to train the model, while the testing set is used to evaluate its performance on unseen data.

```
[26]: # Here we split the data into 80% for training and 20% for testing sets.
split_percent = 0.80
split = int(split_percent * len(X))
```

```
# The first 'split' samples are used for training, and the remaining samples
↪are used for testing.
X_train = X[:split] # Training data for inputs
y_train = y[:split] # Training data for outputs
X_test = X[split:] # Testing data for inputs
y_test = y[split:] # Testing data for outputs
```

### 1.3.4 Building the LSTM Model

In this section, we define the architecture of the LSTM model, compile it, and prepare it for training. The LSTM model is designed to capture temporal dependencies in the data, making it suitable for time series forecasting.

```
[27]: from keras.models import Sequential
      from keras.layers import LSTM, Dense, Dropout

      # Here we add an LSTM layer with 50 units, set return_sequences=True to stack
      ↪another LSTM layer
      lstm_model = Sequential()
      lstm_model.add(LSTM(50, return_sequences=True, input_shape=(look_back, 1)))

      # And here we add another LSTM layer with 50 units
      # This layer will process the output from the previous LSTM layer.
      lstm_model.add(LSTM(50))

      # And here a Dense layer with a single unit used to output the predicted value
      # This is the final output layer that produces the prediction.
      lstm_model.add(Dense(1))

      # The Adam optimizer is used for efficient gradient descent, and mean squared
      ↪error is used as the loss function to minimize prediction error.
      lstm_model.compile(optimizer='adam', loss='mean_squared_error')
```

```
2024-05-28 00:56:25.905230: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
```

```
To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 FMA, in other
operations, rebuild TensorFlow with the appropriate compiler flags.
```

```
2024-05-28 00:56:28.522801: I
tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool
with default inter op setting: 2. Tune using inter_op_parallelism_threads for
best performance.
```

```
2024-05-28 00:56:28.943085: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with
dtype int32
```

```
[[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]
```

```
2024-05-28 00:56:28.944928: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_grad/concat/split/split_dim' with dtype
int32
```

```
[[{{node gradients/split_grad/concat/split/split_dim}}]]
2024-05-28 00:56:28.946538: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with
dtype int32
```

```
[[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]
2024-05-28 00:56:29.278899: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with
dtype int32
```

```
[[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]
2024-05-28 00:56:29.280919: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_grad/concat/split/split_dim' with dtype
int32
```

```
[[{{node gradients/split_grad/concat/split/split_dim}}]]
2024-05-28 00:56:29.282915: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with
dtype int32
```

```
[[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]
```

### 1.3.5 Training the LSTM Model

In this section, we train the LSTM model using the training data. We specify the number of epochs, batch size, and validation data to monitor the model's performance on unseen data during training.

```
[28]: # history = lstm_model.fit(X_train, y_train, epochs=100, batch_size=32,
      ↪ validation_data=(X_test, y_test), verbose=1)
# train_loss = history.history['loss']
# val_loss = history.history['val_loss']

from keras.callbacks import EarlyStopping

# Initialize EarlyStopping to monitor 'val_loss' with patience of 40 epochs
# EarlyStopping helps to prevent overfitting by monitoring the 'val_loss'
  ↪ (validation loss) during training.
```

```

# If the 'val_loss' does not improve for a specified number of epochs
↳(patience), training stops early.
# The 'restore_best_weights' parameter ensures that the model reverts to the
↳weights with the best 'val_loss'.
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=40,
    verbose=1,
    restore_best_weights=True
)

# And here we train the LSTM model with early stopping to prevent overfitting
history = lstm_model.fit(
    X_train, y_train,
    validation_data=(X_test, y_test),
    epochs=110,
    batch_size=64,
    verbose=1,
    callbacks=[early_stopping]
)

```

Epoch 1/100

2024-05-28 00:56:29.876702: I tensorflow/core/common\_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_2\_grad/concat/split\_2/split\_dim' with dtype int32

[[{{node gradients/split\_2\_grad/concat/split\_2/split\_dim}}]]

2024-05-28 00:56:29.879038: I tensorflow/core/common\_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_grad/concat/split/split\_dim' with dtype int32

[[{{node gradients/split\_grad/concat/split/split\_dim}}]]

2024-05-28 00:56:29.881070: I tensorflow/core/common\_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_1\_grad/concat/split\_1/split\_dim' with dtype int32

[[{{node gradients/split\_1\_grad/concat/split\_1/split\_dim}}]]

2024-05-28 00:56:30.238628: I tensorflow/core/common\_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_2\_grad/concat/split\_2/split\_dim' with dtype int32

[[{{node gradients/split\_2\_grad/concat/split\_2/split\_dim}}]]

2024-05-28 00:56:30.241130: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_grad/concat/split/split\_dim' with dtype int32

[[{{node gradients/split\_grad/concat/split/split\_dim}}]]

2024-05-28 00:56:30.243461: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_1\_grad/concat/split\_1/split\_dim' with dtype int32

[[{{node gradients/split\_1\_grad/concat/split\_1/split\_dim}}]]

2024-05-28 00:56:31.877622: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_2\_grad/concat/split\_2/split\_dim' with dtype int32

[[{{node gradients/split\_2\_grad/concat/split\_2/split\_dim}}]]

2024-05-28 00:56:31.880585: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_grad/concat/split/split\_dim' with dtype int32

[[{{node gradients/split\_grad/concat/split/split\_dim}}]]

2024-05-28 00:56:31.883146: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_1\_grad/concat/split\_1/split\_dim' with dtype int32

[[{{node gradients/split\_1\_grad/concat/split\_1/split\_dim}}]]

2024-05-28 00:56:32.247646: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_2\_grad/concat/split\_2/split\_dim' with dtype int32

[[{{node gradients/split\_2\_grad/concat/split\_2/split\_dim}}]]

2024-05-28 00:56:32.249739: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_grad/concat/split/split\_dim' with dtype int32

[[{{node gradients/split\_grad/concat/split/split\_dim}}]]

2024-05-28 00:56:32.251751: I tensorflow/core/common\_runtime/executor.cc:1197]

[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_1\_grad/concat/split\_1/split\_dim' with dtype int32

[[{{node gradients/split\_1\_grad/concat/split\_1/split\_dim}}]]

```

183/183 [=====] - ETA: 0s - loss: 0.0010

2024-05-28 00:56:49.379607: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with
dtype int32
    [[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]
2024-05-28 00:56:49.381585: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_grad/concat/split/split_dim' with dtype
int32
    [[{{node gradients/split_grad/concat/split/split_dim}}]]
2024-05-28 00:56:49.383205: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with
dtype int32
    [[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]
2024-05-28 00:56:49.683670: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with
dtype int32
    [[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]
2024-05-28 00:56:49.685792: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_grad/concat/split/split_dim' with dtype
int32
    [[{{node gradients/split_grad/concat/split/split_dim}}]]
2024-05-28 00:56:49.687506: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with
dtype int32
    [[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]

183/183 [=====] - 22s 86ms/step - loss: 0.0010 -
val_loss: 0.0013
Epoch 2/100
183/183 [=====] - 14s 79ms/step - loss: 4.5341e-05 -
val_loss: 8.3862e-04
Epoch 3/100
183/183 [=====] - 15s 81ms/step - loss: 4.4514e-05 -
val_loss: 8.6462e-04
Epoch 4/100
183/183 [=====] - 13s 72ms/step - loss: 4.2607e-05 -

```

```

val_loss: 5.9635e-04
Epoch 5/100
183/183 [=====] - 14s 76ms/step - loss: 4.4565e-05 -
val_loss: 4.7125e-04
Epoch 6/100
183/183 [=====] - 14s 76ms/step - loss: 4.3605e-05 -
val_loss: 5.9436e-04
Epoch 7/100
183/183 [=====] - 15s 79ms/step - loss: 3.5813e-05 -
val_loss: 5.1969e-04
Epoch 8/100
183/183 [=====] - 14s 79ms/step - loss: 3.6504e-05 -
val_loss: 2.8483e-04
Epoch 9/100
183/183 [=====] - 37s 202ms/step - loss: 3.8940e-05 -
val_loss: 2.6909e-04
Epoch 10/100
183/183 [=====] - 42s 228ms/step - loss: 3.3629e-05 -
val_loss: 2.6580e-04
Epoch 11/100
183/183 [=====] - 22s 121ms/step - loss: 2.9985e-05 -
val_loss: 2.6423e-04
Epoch 12/100
183/183 [=====] - 27s 146ms/step - loss: 3.1191e-05 -
val_loss: 2.4488e-04
Epoch 13/100
183/183 [=====] - 31s 166ms/step - loss: 2.9802e-05 -
val_loss: 2.1264e-04
Epoch 14/100
183/183 [=====] - 22s 119ms/step - loss: 3.1531e-05 -
val_loss: 2.5014e-04
Epoch 15/100
183/183 [=====] - 30s 164ms/step - loss: 2.8835e-05 -
val_loss: 1.9909e-04
Epoch 16/100
183/183 [=====] - 47s 255ms/step - loss: 2.5686e-05 -
val_loss: 2.3386e-04
Epoch 17/100
183/183 [=====] - 23s 125ms/step - loss: 2.5330e-05 -
val_loss: 2.1470e-04
Epoch 18/100
183/183 [=====] - 26s 144ms/step - loss: 2.6641e-05 -
val_loss: 1.8590e-04
Epoch 19/100
183/183 [=====] - 24s 129ms/step - loss: 2.2585e-05 -
val_loss: 2.3322e-04
Epoch 20/100
183/183 [=====] - 25s 135ms/step - loss: 2.2476e-05 -

```

```

val_loss: 1.7421e-04
Epoch 21/100
183/183 [=====] - 22s 119ms/step - loss: 2.2802e-05 -
val_loss: 1.7330e-04
Epoch 22/100
183/183 [=====] - 22s 121ms/step - loss: 2.3199e-05 -
val_loss: 1.5303e-04
Epoch 23/100
183/183 [=====] - 21s 112ms/step - loss: 1.9850e-05 -
val_loss: 1.6332e-04
Epoch 24/100
183/183 [=====] - 24s 129ms/step - loss: 2.0309e-05 -
val_loss: 1.4530e-04
Epoch 25/100
183/183 [=====] - 23s 123ms/step - loss: 1.8758e-05 -
val_loss: 1.4957e-04
Epoch 26/100
183/183 [=====] - 22s 118ms/step - loss: 1.8138e-05 -
val_loss: 1.2313e-04
Epoch 27/100
183/183 [=====] - 28s 152ms/step - loss: 1.7423e-05 -
val_loss: 2.3637e-04
Epoch 28/100
183/183 [=====] - 22s 118ms/step - loss: 1.7379e-05 -
val_loss: 1.4354e-04
Epoch 29/100
183/183 [=====] - 20s 109ms/step - loss: 1.6357e-05 -
val_loss: 1.5540e-04
Epoch 30/100
183/183 [=====] - 21s 116ms/step - loss: 1.6346e-05 -
val_loss: 1.7793e-04
Epoch 31/100
183/183 [=====] - 23s 128ms/step - loss: 1.6197e-05 -
val_loss: 1.1832e-04
Epoch 32/100
183/183 [=====] - 22s 119ms/step - loss: 1.4858e-05 -
val_loss: 1.3972e-04
Epoch 33/100
183/183 [=====] - 24s 132ms/step - loss: 1.6183e-05 -
val_loss: 1.1580e-04
Epoch 34/100
183/183 [=====] - 21s 112ms/step - loss: 1.5125e-05 -
val_loss: 1.3607e-04
Epoch 35/100
183/183 [=====] - 22s 122ms/step - loss: 1.4791e-05 -
val_loss: 1.6416e-04
Epoch 36/100
183/183 [=====] - 23s 126ms/step - loss: 1.7144e-05 -

```



```

val_loss: 4.5357e-04
Epoch 37/100
183/183 [=====] - 20s 110ms/step - loss: 1.4584e-05 -
val_loss: 1.1056e-04
Epoch 38/100
183/183 [=====] - 23s 124ms/step - loss: 1.2618e-05 -
val_loss: 1.0436e-04
Epoch 39/100
183/183 [=====] - 23s 128ms/step - loss: 1.3084e-05 -
val_loss: 2.0831e-04
Epoch 40/100
183/183 [=====] - 31s 171ms/step - loss: 1.2777e-05 -
val_loss: 2.5897e-04
Epoch 41/100
183/183 [=====] - 24s 134ms/step - loss: 1.3124e-05 -
val_loss: 2.5948e-04
Epoch 42/100
183/183 [=====] - 34s 187ms/step - loss: 1.2693e-05 -
val_loss: 1.2335e-04
Epoch 43/100
183/183 [=====] - 20s 109ms/step - loss: 1.3156e-05 -
val_loss: 2.2848e-04
Epoch 44/100
183/183 [=====] - 14s 77ms/step - loss: 1.2919e-05 -
val_loss: 1.4552e-04
Epoch 45/100
183/183 [=====] - 13s 71ms/step - loss: 1.2720e-05 -
val_loss: 2.3227e-04
Epoch 46/100
183/183 [=====] - 13s 72ms/step - loss: 1.2804e-05 -
val_loss: 2.2093e-04
Epoch 47/100
183/183 [=====] - 15s 81ms/step - loss: 1.2899e-05 -
val_loss: 2.9371e-04
Epoch 48/100
183/183 [=====] - 15s 80ms/step - loss: 1.3754e-05 -
val_loss: 1.6973e-04
Epoch 49/100
183/183 [=====] - 14s 74ms/step - loss: 1.2489e-05 -
val_loss: 1.6787e-04
Epoch 50/100
183/183 [=====] - 14s 76ms/step - loss: 1.2328e-05 -
val_loss: 2.0934e-04
Epoch 51/100
183/183 [=====] - 14s 75ms/step - loss: 1.1955e-05 -
val_loss: 2.4619e-04
Epoch 52/100
183/183 [=====] - 13s 70ms/step - loss: 1.2549e-05 -

```

```

val_loss: 2.9444e-04
Epoch 53/100
183/183 [=====] - 13s 73ms/step - loss: 1.2357e-05 -
val_loss: 5.0619e-04
Epoch 54/100
183/183 [=====] - 14s 74ms/step - loss: 1.1702e-05 -
val_loss: 2.5570e-04
Epoch 55/100
183/183 [=====] - 14s 76ms/step - loss: 1.2112e-05 -
val_loss: 3.0766e-04
Epoch 56/100
183/183 [=====] - 12s 67ms/step - loss: 1.2254e-05 -
val_loss: 3.9029e-04
Epoch 57/100
183/183 [=====] - 13s 69ms/step - loss: 1.2571e-05 -
val_loss: 2.0947e-04
Epoch 58/100
183/183 [=====] - 14s 77ms/step - loss: 1.3028e-05 -
val_loss: 4.3470e-04
Epoch 59/100
183/183 [=====] - 14s 76ms/step - loss: 1.2106e-05 -
val_loss: 3.4239e-04
Epoch 60/100
183/183 [=====] - 14s 77ms/step - loss: 1.1345e-05 -
val_loss: 3.7403e-04
Epoch 61/100
183/183 [=====] - 13s 73ms/step - loss: 1.2305e-05 -
val_loss: 2.4002e-04
Epoch 62/100
183/183 [=====] - 14s 78ms/step - loss: 1.2092e-05 -
val_loss: 2.3766e-04
Epoch 63/100
183/183 [=====] - 13s 68ms/step - loss: 1.1761e-05 -
val_loss: 3.2236e-04
Epoch 64/100
183/183 [=====] - 14s 77ms/step - loss: 1.1708e-05 -
val_loss: 3.4962e-04
Epoch 65/100
183/183 [=====] - 14s 77ms/step - loss: 1.3452e-05 -
val_loss: 3.3450e-04
Epoch 66/100
183/183 [=====] - 14s 77ms/step - loss: 1.1267e-05 -
val_loss: 4.4552e-04
Epoch 67/100
183/183 [=====] - 13s 73ms/step - loss: 1.2208e-05 -
val_loss: 5.6792e-04
Epoch 68/100
183/183 [=====] - 12s 68ms/step - loss: 1.2589e-05 -

```

```

val_loss: 3.2028e-04
Epoch 69/100
183/183 [=====] - 13s 73ms/step - loss: 1.2110e-05 -
val_loss: 4.7077e-04
Epoch 70/100
183/183 [=====] - 14s 76ms/step - loss: 1.2196e-05 -
val_loss: 3.6355e-04
Epoch 71/100
183/183 [=====] - 13s 69ms/step - loss: 1.1708e-05 -
val_loss: 3.9769e-04
Epoch 72/100
183/183 [=====] - 15s 85ms/step - loss: 1.1407e-05 -
val_loss: 5.1830e-04
Epoch 73/100
183/183 [=====] - 14s 75ms/step - loss: 1.2424e-05 -
val_loss: 5.6920e-04
Epoch 74/100
183/183 [=====] - 14s 77ms/step - loss: 1.1652e-05 -
val_loss: 6.1320e-04
Epoch 75/100
183/183 [=====] - 13s 71ms/step - loss: 1.2472e-05 -
val_loss: 4.0473e-04
Epoch 76/100
183/183 [=====] - 14s 77ms/step - loss: 1.1875e-05 -
val_loss: 5.5109e-04
Epoch 77/100
183/183 [=====] - 15s 80ms/step - loss: 1.2658e-05 -
val_loss: 6.2655e-04
Epoch 78/100
183/183 [=====] - 13s 73ms/step - loss: 1.1803e-05 -
val_loss: 4.8450e-04
Epoch 79/100
183/183 [=====] - 15s 80ms/step - loss: 1.2501e-05 -
val_loss: 5.2640e-04
Epoch 80/100
183/183 [=====] - 13s 70ms/step - loss: 1.2373e-05 -
val_loss: 4.9429e-04
Epoch 81/100
183/183 [=====] - 14s 75ms/step - loss: 1.1757e-05 -
val_loss: 5.3522e-04
Epoch 82/100
183/183 [=====] - 25503s 140s/step - loss: 1.2580e-05 -
val_loss: 3.1446e-04
Epoch 83/100
183/183 [=====] - 35s 188ms/step - loss: 1.1554e-05 -
val_loss: 4.0986e-04
Epoch 84/100
183/183 [=====] - 26s 143ms/step - loss: 1.2342e-05 -

```

```

val_loss: 4.1288e-04
Epoch 85/100
183/183 [=====] - 18s 99ms/step - loss: 1.0869e-05 -
val_loss: 4.7506e-04
Epoch 86/100
183/183 [=====] - 17s 92ms/step - loss: 1.1048e-05 -
val_loss: 4.6264e-04
Epoch 87/100
183/183 [=====] - 17s 95ms/step - loss: 1.1666e-05 -
val_loss: 4.7030e-04
Epoch 88/100
183/183 [=====] - 19s 102ms/step - loss: 1.1238e-05 -
val_loss: 6.4671e-04
Epoch 89/100
183/183 [=====] - 15s 82ms/step - loss: 1.1501e-05 -
val_loss: 5.3281e-04
Epoch 90/100
183/183 [=====] - 17s 96ms/step - loss: 1.1278e-05 -
val_loss: 6.5807e-04
Epoch 91/100
183/183 [=====] - 15s 82ms/step - loss: 1.1165e-05 -
val_loss: 5.3917e-04
Epoch 92/100
183/183 [=====] - 18s 101ms/step - loss: 1.2720e-05 -
val_loss: 3.4536e-04
Epoch 93/100
183/183 [=====] - 14s 78ms/step - loss: 1.0896e-05 -
val_loss: 2.9224e-04
Epoch 94/100
183/183 [=====] - 16s 86ms/step - loss: 1.1699e-05 -
val_loss: 3.3770e-04
Epoch 95/100
183/183 [=====] - 15s 80ms/step - loss: 1.3322e-05 -
val_loss: 4.0450e-04
Epoch 96/100
183/183 [=====] - 13s 72ms/step - loss: 1.1599e-05 -
val_loss: 3.3362e-04
Epoch 97/100
183/183 [=====] - 16s 89ms/step - loss: 1.1108e-05 -
val_loss: 7.5887e-04
Epoch 98/100
183/183 [=====] - 17s 93ms/step - loss: 1.1100e-05 -
val_loss: 4.8278e-04
Epoch 99/100
183/183 [=====] - 16s 85ms/step - loss: 1.1414e-05 -
val_loss: 4.3803e-04
Epoch 100/100
183/183 [=====] - 17s 95ms/step - loss: 1.1957e-05 -

```

val\_loss: 2.9018e-04

### 1.3.6 Making Predictions and Inverting the Scale

In this section, we make predictions using the trained LSTM model and then invert the predictions back to the original scale to interpret the results correctly.

```
[29]: # Here we make predictions with the trained LSTM model for both training and
      ↪testing data
y_train_lstm_pred = lstm_model.predict(X_train)
y_test_lstm_pred = lstm_model.predict(X_test)

# Invert predictions back to original scale
# This step is necessary because the data was scaled before training.
# The scaler.inverse_transform() function is used to convert the scaled values
↪back to their original scale.
y_train_lstm_pred = scaler.inverse_transform(y_train_lstm_pred)
y_train = scaler.inverse_transform([y_train])
y_test_lstm_pred = scaler.inverse_transform(y_test_lstm_pred)
y_test = scaler.inverse_transform([y_test])
```

```
2024-05-28 08:32:24.419350: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with
dtype int32
```

```
[[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]
2024-05-28 08:32:24.421438: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_grad/concat/split/split_dim' with dtype
int32
```

```
[[{{node gradients/split_grad/concat/split/split_dim}}]]
2024-05-28 08:32:24.424167: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with
dtype int32
```

```
[[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]
2024-05-28 08:32:24.809031: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with
dtype int32
```

```
[[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]
2024-05-28 08:32:24.812176: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
```

error and you can ignore this message): INVALID\_ARGUMENT: You must feed a value for placeholder tensor 'gradients/split\_grad/concat/split/split\_dim' with dtype int32

```
[[{{node gradients/split_grad/concat/split/split_dim}}]]
2024-05-28 08:32:24.814757: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with
dtype int32
```

```
[[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]
```

```
183/183 [=====] - 8s 38ms/step
```

```
46/46 [=====] - 2s 47ms/step
```

### 1.3.7 Plotting Training and Validation Loss

In this section, we visualize the training and validation loss over the epochs to evaluate the model's performance and identify the point of minimum validation loss.

```
[30]: # import matplotlib.pyplot as plt
# import numpy as np

# # Assuming 'history.history['val_loss']' contains your validation loss data
# val_loss = np.array(history.history['val_loss'])
# min_val_loss_idx = np.argmin(val_loss) # Get the index of the minimum
# validation loss
# min_val_loss = val_loss[min_val_loss_idx] # Get the minimum validation loss
# value

# plt.figure(figsize=(10, 5))
# plt.plot(history.history['loss'], label='Training Loss')
# plt.plot(history.history['val_loss'], label='Validation Loss')
# plt.scatter(min_val_loss_idx, min_val_loss, color='red') # Mark the min point

# # Annotate the min point automatically
# plt.annotate(f'{min_val_loss_idx + 1}',
#             xy=(min_val_loss_idx, min_val_loss),
#             xytext=(min_val_loss_idx, min_val_loss - 0.0001), # Position it
#             # slightly below the point
#             textcoords='offset points',
#             arrowprops=dict(arrowstyle='->', color='red'),
#             ha='center', va='top', color='red', fontsize=12,
#             # fontweight='bold')

# plt.xticks(fontsize=14) # Larger font size for the x-axis ticks
# plt.yticks(fontsize=14)

# # Add labels and title
```

```
# plt.xlabel('Epochs', fontsize=18)
# plt.ylabel('Loss', fontsize=18)
# plt.title('Model Training and Validation Loss', fontsize=18)
# plt.legend(prop={'size': 18, 'weight': 'bold'})
# plt.savefig('ep_graph1.pdf', format='pdf')
# plt.show()
```

### 1.3.8 Organizing and Inspecting LSTM Prediction Results in Table Form

In this section, we consolidate the predictions with the actual values into structured DataFrames, aligning them with their corresponding dates. This organization is essential for an intuitive inspection of the model's predictive accuracy. It also lays the groundwork for subsequent analysis, such as calculating error metrics and visualizing the results.

```
[31]: if 'Date' in df.columns:
        df['Date'] = pd.to_datetime(df['Date'])

look_back = 30 # This should be the same look_back you used earlier in your
               ↪ model

# Ensure the date slices start from the correct index
train_dates = df['Date'][look_back:look_back+len(y_train_lstm_pred)].
               ↪ reset_index(drop=True)
test_dates = df['Date'][look_back+len(y_train_lstm_pred):
                       ↪ look_back+len(y_train_lstm_pred)+len(y_test_lstm_pred)].
               ↪ reset_index(drop=True)

# Ensure the close price slices start from the correct index and match the
               ↪ length of the predictions
train_actual_close = df['Close'][look_back:look_back+len(y_train_lstm_pred)].
                    ↪ values
test_actual_close = df['Close'][look_back+len(y_train_lstm_pred):
                               ↪ look_back+len(y_train_lstm_pred)+len(y_test_lstm_pred)].values

# Create the DataFrame using the aligned data
lstm_train_results = pd.DataFrame({
    'Date': train_dates,
    'Actual_Close': train_actual_close,
    'Predicted_Close': y_train_lstm_pred.flatten()
})

lstm_test_results = pd.DataFrame({
    'Date': test_dates,
    'Actual_Close': test_actual_close,
    'Predicted_Close': y_test_lstm_pred.flatten()
})
```

```
print("Train Results:\n", lstm_train_results)
print("\nTest Results:\n", lstm_test_results)
```

Train Results:

	Date	Actual_Close	Predicted_Close
0	1995-02-14	482.549988	484.415375
1	1995-02-15	484.540009	485.242249
2	1995-02-16	485.220001	486.972473
3	1995-02-17	481.970001	487.840210
4	1995-02-21	482.720001	485.512848
...	...	...	...
5843	2018-05-01	2654.800049	2660.910645
5844	2018-05-02	2635.669922	2670.742432
5845	2018-05-03	2629.729980	2647.377197
5846	2018-05-04	2663.419922	2644.906006
5847	2018-05-07	2672.629883	2680.181152

[5848 rows x 3 columns]

Test Results:

	Date	Actual_Close	Predicted_Close
0	2018-05-08	2671.919922	2685.506592
1	2018-05-09	2697.790039	2685.966797
2	2018-05-10	2723.070068	2713.422607
3	2018-05-11	2727.719971	2736.490723
4	2018-05-14	2730.129883	2739.522461
...	...	...	...
1457	2024-02-22	5087.029785	4802.567871
1458	2024-02-23	5088.799805	4903.547363
1459	2024-02-26	5069.529785	4877.709473
1460	2024-02-27	5078.180176	4874.244141
1461	2024-02-28	5069.759766	4879.318359

[1462 rows x 3 columns]

### 1.3.9 Performance Metrics Evaluation for LSTM

In this segment, we compute and display the performance metrics for both the training and testing datasets. This evaluation involves Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE) and Root Mean Squared Error (RMSE). These metrics help to quantify the accuracy of our model and reveal how well the predictions match up with the actual stock prices.

```
[32]: from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

# Function to calculate MAPE
def mean_absolute_percentage_error(y_true, y_pred):
```



```

    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

# Calculate metrics for the training set
lstm_train_mae = mean_absolute_error(lstm_train_results['Actual_Close'],
    ↪lstm_train_results['Predicted_Close'])
lstm_train_rmse = np.
    ↪sqrt(mean_squared_error(lstm_train_results['Actual_Close'],
    ↪lstm_train_results['Predicted_Close']))
lstm_train_mape =
    ↪mean_absolute_percentage_error(lstm_train_results['Actual_Close'],
    ↪lstm_train_results['Predicted_Close'])

# Calculate metrics for the testing set
lstm_test_mae = mean_absolute_error(lstm_test_results['Actual_Close'],
    ↪lstm_test_results['Predicted_Close'])
lstm_test_rmse = np.sqrt(mean_squared_error(lstm_test_results['Actual_Close'],
    ↪lstm_test_results['Predicted_Close']))
lstm_test_mape =
    ↪mean_absolute_percentage_error(lstm_test_results['Actual_Close'],
    ↪lstm_test_results['Predicted_Close'])

# Print out the metrics for the training set
print("Training set metrics:")
print(f"Mean Absolute Error (MAE): {lstm_train_mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {lstm_train_rmse:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {lstm_train_mape:.2f}%")

# Print out the metrics for the testing set
print("\nTesting set metrics:")
print(f"Mean Absolute Error (MAE): {lstm_test_mae:.2f}")
print(f"Root Mean Squared Error (RMSE): {lstm_test_rmse:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {lstm_test_mape:.2f}%")

```

```

Training set metrics:
Mean Absolute Error (MAE): 10.89
Root Mean Squared Error (RMSE): 15.57
Mean Absolute Percentage Error (MAPE): 0.84%

```

```

Testing set metrics:
Mean Absolute Error (MAE): 59.42
Root Mean Squared Error (RMSE): 78.99
Mean Absolute Percentage Error (MAPE): 1.51%

```

### 1.3.10 Visualizing LSTM Model Predictions Against Actual Data

In this section, we visualize the predicted stock prices from the LSTM model against the actual S&P 500 closing prices. This step provides a visual representation of the model's performance over time, showcasing how well the predictions align with real-world data.

```
[33]: import matplotlib.pyplot as plt
import pandas as pd

# Combine train and test results into a single DataFrame
lstm_combined_results = pd.concat([lstm_train_results, lstm_test_results])

# Convert 'Date' to datetime and sort by date to ensure correct plotting order
lstm_combined_results['Date'] = pd.to_datetime(lstm_combined_results['Date'])
lstm_combined_results.sort_values('Date', inplace=True)

# Set 'Date' as the index for plotting
lstm_combined_results.set_index('Date', inplace=True)


# Plot the actual close prices
plt.figure(figsize=(14,7))
plt.plot(lstm_combined_results['Actual_Close'], label='Actual Close',
        color='blue', linestyle='--')

# Plot the training predictions
plt.plot(lstm_train_results['Date'], lstm_train_results['Predicted_Close'],
        label='Training Predictions', color='orange')

# Plot the testing predictions
plt.plot(lstm_test_results['Date'], lstm_test_results['Predicted_Close'],
        label='Testing Predictions', color='green')

plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

# Added labels and title
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price', fontsize=18)
plt.title('Actual Close vs Predictions', fontsize=18)
plt.legend(prop={'size': 18, 'weight': 'bold'})
plt.savefig('lstm_result_graph.pdf', format='pdf')
plt.show()
```



## 1.4 Comparison for All\_Models together

### 1.4.1 Creating a Comparison Table for All Models

In this section, we prepare and merge the prediction results from the polynomial regression, decision tree, and LSTM models into a single table for easy comparison.

```
[34]: # Convert 'Date' columns to datetime if they're not already
poly_test_results['Date'] = pd.to_datetime(poly_test_results['Date'])
decision_test_results['Date'] = pd.to_datetime(decision_test_results['Date'])
lstm_test_results['Date'] = pd.to_datetime(lstm_test_results['Date'])

poly_train_results['Date'] = pd.to_datetime(poly_train_results['Date'])
decision_train_results['Date'] = pd.to_datetime(decision_train_results['Date'])
lstm_train_results['Date'] = pd.to_datetime(lstm_train_results['Date'])

# Merge the train and test results for each model
poly_results = pd.concat([poly_train_results, poly_test_results]).
    ↪sort_values('Date')
decision_results = pd.concat([decision_train_results, decision_test_results]).
    ↪sort_values('Date')
lstm_results = pd.concat([lstm_train_results, lstm_test_results]).
    ↪sort_values('Date')

# Merge the results on the Date field
merged_results = pd.merge(poly_results, decision_results, on=['Date',
    ↪'Actual_Close'])
```

```

combined_results = pd.merge(merged_results, lstm_results, on=['Date',
↪ 'Actual_Close'])
#merged_resultss = pd.merge(poly_results, decision_results, lstm_results,
↪ on=['Date', 'Actual_Close'])

# Rename columns for clarity
combined_results.rename(columns={
    'Poly_Predicted_Close': 'Poly_Predicted',
    'Predicted_Close_x': 'Decision_Tree_Predicted',
    'Predicted_Close_y': 'LSTM_Predicted'
}, inplace=True)

# Set the 'Date' column as index
combined_results.set_index('Date', inplace=True)

# print(combined_results.to_string())

print("Train Results:\n", combined_results.head().to_string())
print("\nTest Results:\n", combined_results.tail().to_string())

```

Train Results:

	Actual_Close	Poly_Predicted	Decision_Tree_Predicted
LSTM_Predicted			
Date			
1995-02-14	482.549988	679.072165	486.522944
484.415375			
1995-02-15	484.540009	690.029350	486.522944
485.242249			
1995-02-16	485.220001	686.626128	486.522944
486.972473			
1995-02-17	481.970001	685.117534	486.522944
487.840210			
1995-02-21	482.720001	679.872469	486.522944
485.512848			

Test Results:

	Actual_Close	Poly_Predicted	Decision_Tree_Predicted
LSTM_Predicted			
Date			
2024-02-22	5087.029785	8765.893054	2808.062512
4802.567871			
2024-02-23	5088.799805	8617.675638	2808.062512
4903.547363			
2024-02-26	5069.529785	8548.768183	2808.062512
4877.709473			
2024-02-27	5078.180176	8455.260613	2808.062512

4874.244141  
2024-02-28    5069.759766        8470.372507                    2808.062512  
4879.318359

### 1.4.2 Plotting Actual Close vs. Predictions for All Models


In this section, we plot the actual closing prices against the predictions made by the polynomial regression, decision tree, and LSTM models to visually compare their performances.

```
[35]: import pandas as pd
import matplotlib.pyplot as plt

plt.figure(figsize=(14, 7))
plt.plot(combined_results.index, combined_results['Actual_Close'],
         ↪label='Actual Close', color='blue', linestyle='-')
plt.plot(combined_results.index, combined_results['Poly_Predicted'],
         ↪label='Polynomial Predictions', color='green', linestyle='-')
plt.plot(combined_results.index, combined_results['Decision_Tree_Predicted'],
         ↪label='Decision Tree Predictions', color='orange', linestyle='-')
plt.plot(combined_results.index, combined_results['LSTM_Predicted'],
         ↪label='LSTM Predictions', color='red', linestyle='-')

plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price', fontsize=18)
plt.title('Actual Close vs Predictions', fontsize=18)
plt.legend(prop={'size': 18, 'weight': 'bold'})
plt.savefig('actual_vs_predictions_graph1.pdf', format='pdf')
plt.show()
```



### 1.4.3 Zoomed-In Plot of Actual Close vs. Model Predictions (Testing Phase)

In this section, we create a zoomed-in plot to focus on the testing results, showing the performance of polynomial regression, decision tree, and LSTM models from 2018 onwards.

```
[36]: import pandas as pd
import matplotlib.pyplot as plt

if 'Date' not in combined_results.columns:
    combined_results.reset_index(inplace=True)

combined_results['Date'] = pd.to_datetime(combined_results['Date'])


combined_results.set_index('Date', inplace=True)

zoomed_results = combined_results['2018-01-01:']

plt.figure(figsize=(16, 10))
plt.plot(zoomed_results['Actual_Close'], label='Actual Close', color='blue',
        linestyle='-')
plt.plot(zoomed_results['Poly_Predicted'], label='Polynomial Predictions',
        color='green', linestyle='-')
plt.plot(zoomed_results['Decision_Tree_Predicted'], label='Decision Tree
        Predictions', color='orange', linestyle='-')
```

```
plt.plot(zoomed_results['LSTM_Predicted'], label='LSTM Predictions',
        color='red', linestyle='-')

plt.xticks(fontsize=14) # Larger font size for the x-axis ticks
plt.yticks(fontsize=14)
# Add labels and title
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price', fontsize=18)
plt.title('Actual Close vs Model Predictions (2015 Onwards)', fontsize=18)
plt.legend(prop={'size': 18, 'weight': 'bold'})
plt.savefig('zoomed_test_predictions_graph3.pdf', format='pdf')
plt.show()
```



#### 1.4.4 Zoomed-In Plot of Actual Close vs. Model Predictions (Training Phase)

In this section, we create a zoomed-in plot to focus on the training results, showing the performance of polynomial regression, decision tree, and LSTM models from 2004 to 2007.

```
[37]: import pandas as pd
import matplotlib.pyplot as plt

# Ensure 'Date' is a column
if 'Date' not in combined_results.columns:
```

```

combined_results.reset_index(inplace=True)

combined_results['Date'] = pd.to_datetime(combined_results['Date'])

# Set the 'Date' column as index
combined_results.set_index('Date', inplace=True)


zoomed_results = combined_results['2004-08-01':'2006-12-31']

# Plotting
plt.figure(figsize=(16, 10))
plt.plot(zoomed_results['Actual_Close'], label='Actual Close', color='blue',
        ↪linestyle='-')
plt.plot(zoomed_results['Poly_Predicted'], label='Polynomial Predictions',
        ↪color='green', linestyle='-')
plt.plot(zoomed_results['Decision_Tree_Predicted'], label='Decision Tree
        ↪Predictions', color='orange', linestyle='-')
plt.plot(zoomed_results['LSTM_Predicted'], label='LSTM Predictions',
        ↪color='red', linestyle='-')

plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
# Add labels and title
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price', fontsize=18)
plt.title('Actual Close vs Model Predictions (2004-2007) in Training phase',
        ↪fontsize=18)
plt.legend(prop={'size': 18, 'weight': 'bold'})
plt.savefig('zoomed_training_predictions_graph2.pdf', format='pdf')
plt.show()

```





### 1.4.5 Comparison of Model Performance Metrics

To compare the performance of different models (Polynomial Regression, Decision Tree, and LSTM), we plot the Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and Mean Absolute Percentage Error (MAPE) for both the training and testing datasets. These metrics provide insights into the accuracy and reliability of each model.

```
[38]: import matplotlib.pyplot as plt
import pandas as pd

# Sample data setup
data = {
    'Model': ['Polynomial', 'Decision Tree', 'LSTM'],
    'MAE_Train': [poly_train_mae, dt_train_mae, lstm_train_mae],
    'RMSE_Train': [poly_train_rmse, dt_train_rmse, lstm_train_rmse],
    'MAPE_Train': [poly_train_mape, dt_train_mape, lstm_train_mape],
    'MAE_Test': [poly_test_mae, dt_test_mae, lstm_test_mae],
    'RMSE_Test': [poly_test_rmse, dt_test_rmse, lstm_test_rmse],
    'MAPE_Test': [poly_test_mape, dt_test_mape, lstm_test_mape],
}

df = pd.DataFrame(data)

def plot_metric(metric):
```

```

fig, ax = plt.subplots(figsize=(8, 6))
metric_train = metric + '_Train'
metric_test = metric + '_Test'
ind = df['Model']
train = df[metric_train]
test = df[metric_test]

# Create bar plot
bars1 = ax.bar(ind, train, width=0.4, label='Train', align='center')
bars2 = ax.bar(ind, test, width=0.4, label='Test', align='edge')

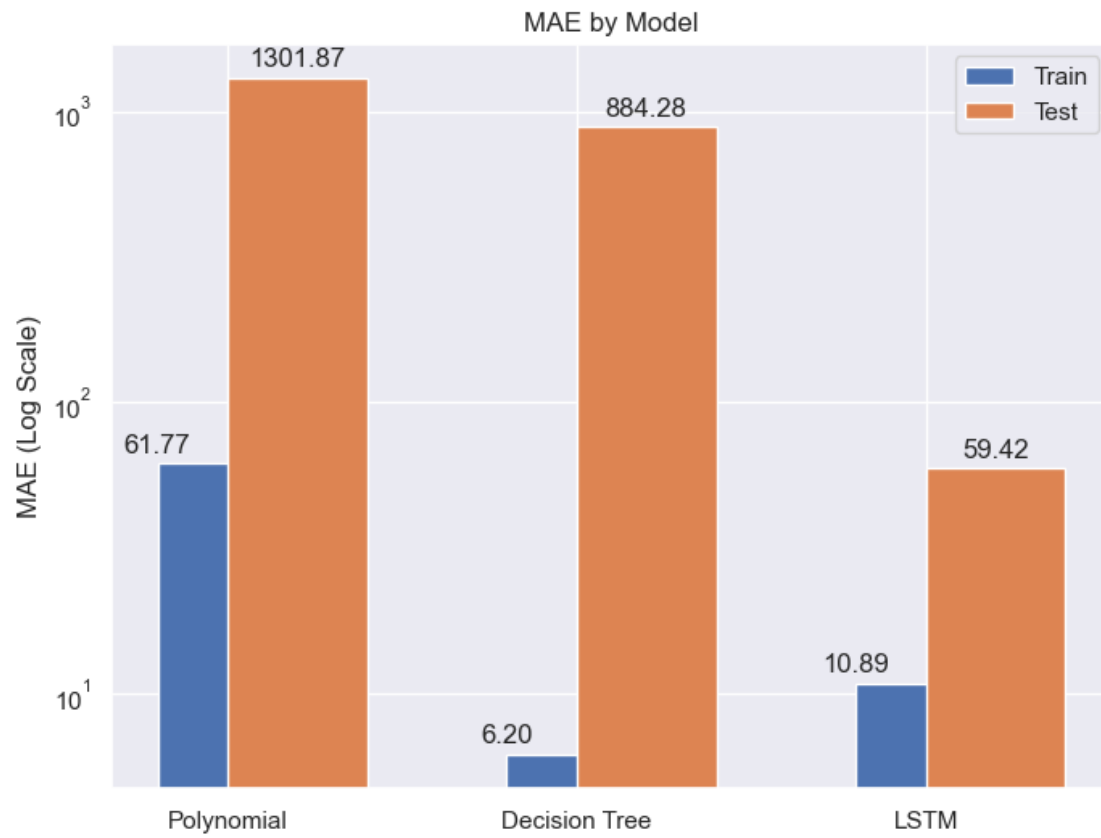
if metric == 'MAPE':
    ax.set_ylabel(f'{metric} (%)') # Set y-label as percentage for MAPE
    for bar in bars1:
        height = bar.get_height()
        ax.annotate(f'{height:.2f}%', xy=(bar.get_x() + bar.get_width() / 2,
↪ 0.2, height),
                    xytext=(0, 3), textcoords='offset points', ha='center',
↪ va='bottom')
    for bar in bars2:
        height = bar.get_height()
        ax.annotate(f'{height:.2f}%', xy=(bar.get_x() + bar.get_width() /
↪ 2, height),
                    xytext=(0, 3), textcoords='offset points', ha='center',
↪ va='bottom')
    else:
        ax.set_ylabel(f'{metric} (Log Scale)')
        ax.set_yscale('log') # Set y-scale to log for MAE and RMSE
        for bar in bars1:
            height = bar.get_height()
            ax.annotate(f'{height:.2f}', xy=(bar.get_x() + bar.get_width() / 2,
↪ 0.2, height),
                        xytext=(0, 3), textcoords='offset points', ha='center',
↪ va='bottom')
        for bar in bars2:
            height = bar.get_height()
            ax.annotate(f'{height:.2f}', xy=(bar.get_x() + bar.get_width() / 2,
↪ height),
                        xytext=(0, 3), textcoords='offset points', ha='center',
↪ va='bottom')


ax.set_title(f'{metric} by Model')
ax.legend()

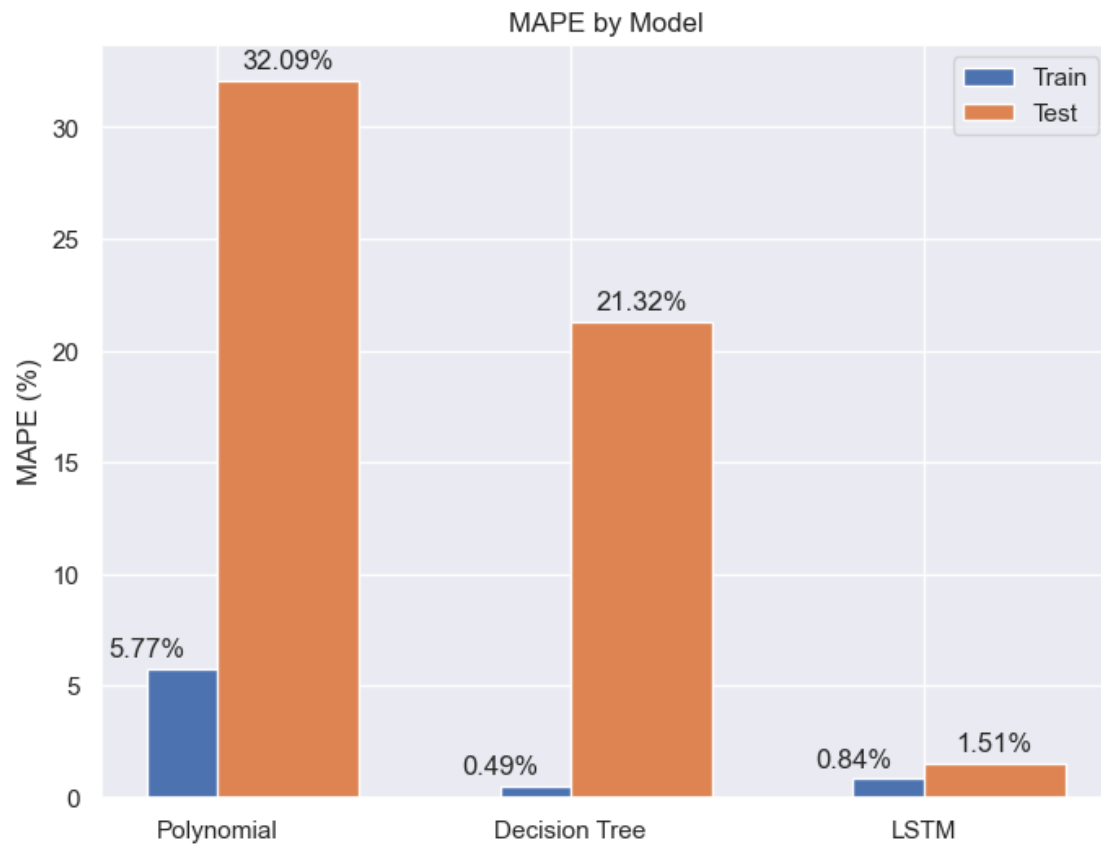
plt.savefig(f'{metric}_comparison_evaluation.pdf', format='pdf')
plt.show()

```

```
# Plotting each metric
metrics = ['MAE', 'RMSE', 'MAPE']
for metric in metrics:
    plot_metric(metric)
```







[ ]:

[ ]:

[ ]: